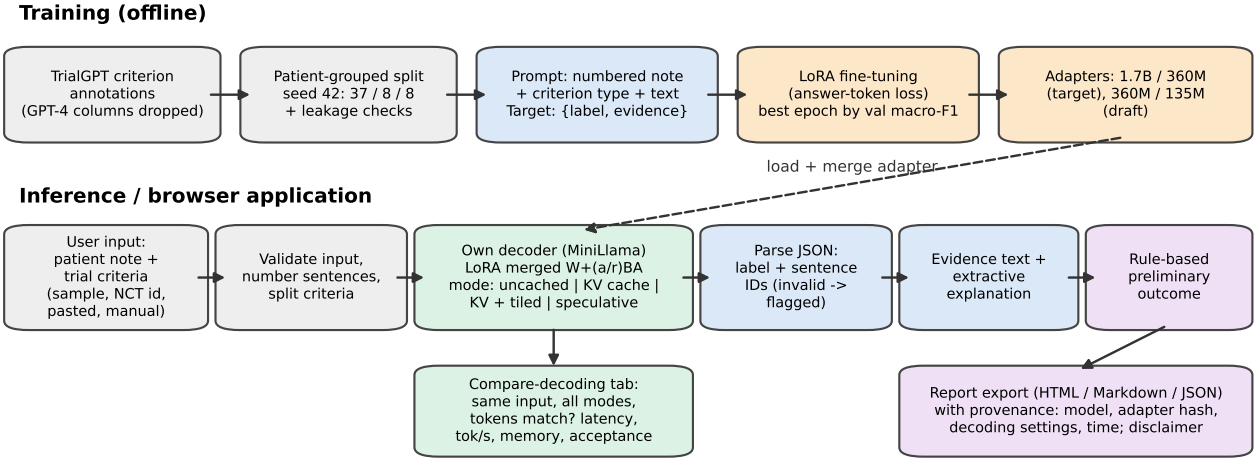
Training (offline)
TrialGPT criterion annotations (GPT-4 columns dropped)
Patient-grouped split seed 42: 37 / 8 / 8 + leakage checks
Prompt: numbered note + criterion type + text Target: {label, evidence}
LoRA fine-tuning (answer-token loss) best epoch by val macro-F1
Adapters: 1.7B / 360M (target), 360M / 135M (draft)
load + merge adapter
Inference / browser application
User input: patient note + trial criteria (sample, NCT id, pasted, manual)
Validate input, number sentences, split criteria
Own decoder (MiniLlama) LoRA merged W+(a/r)BA mode: uncached | KV cache | KV + tiled | speculative
Parse JSON: label + sentence IDs (invalid -> flagged)
Evidence text + extractive explanation
Rule-based preliminary outcome
Compare-decoding tab: same input, all modes, tokens match? latency, tok/s, memory, acceptance
Report export (HTML / Markdown / JSON) with provenance: model, adapter hash, decoding settings, time; disclaimer